TELCO CHURN ANALYSIS

Final Term project

David LOUIS
louisd2025@fau.edu

Abstract

Customer retention is a major challenge for all companies, as acquiring new customers is significantly more expensive than retaining existing ones. Predicting churn helps companies identify customers who are likely to leave and take proactive measures to retain them.

In this project, it begins with an in-depth **Exploratory Data Analysis (EDA)** to uncover the key **Demographic**, **Contractual** and **Service Usage** characteristics that drive **Churn**. Based on this information, we develop two complementary machine learning approaches: **A Logistic Regression Classifier** to predict individual churn risk, and a **K-Means Clustering** model to segment the customer base into high- and low-risk groups.

Together, these models provide both granular predictions and actionable customer segments, enabling the Telecom company to tailor its loyalty strategies and maximize long-term revenues.

Dataset Description

The Telecom Churn Dataset from Kaggle contains **Customer Data**, including ***Demographic Details, Service Usage, Contractual Information, Payment Method...***

The dataset includes the following characteristics:

- Size: 7043 entries
- Type: Tabular data
- Number of Variables: 21 columns
- Target Variable: Churn (indicating whether a customer has left the service)
- Source: https://www.kaggle.com/datasets/blastchar/telco-customer-churn

Introduction

In today's hyper-competitive telecommunications sector, retaining existing customers is both more profitable and strategically more important than acquiring new ones. Average customer acquisition costs can exceed $300 per subscriber, while additional revenues from renewals or additional sales contribute directly to the bottom line. Yet industry-wide churn rates often hover between 15% and 30% per year, resulting in millions of dollars in lost revenue and unnecessary marketing spend. In this context, data-driven churn analysis offers telecom operators a powerful leverage: by identifying which customers are most likely to leave and why, companies can then proactively adapt offers, optimize service offerings and deploy targeted loyalty campaigns to reduce churn.

Business Context

Our client, a mid-sized Telecom provider, faces rising customer turnover driven by flexible month-to-month contracts, aggressive pricing wars on fiber-optic plans, and fragmented service

offerings. As competitors bundle streaming, security, and device-protection packages, customers increasingly shop around, weighing both cost and service breadth. The ability to pinpoint high-risk subscribers before they defect and to understand the key factors behind their decisions is essential to stabilizing total revenue, improving customer lifetime value, and maintaining competitive differentiation in a saturated market.

Project Objectives

1. **Exploratory Data Analysis (EDA):**
    a. Uncover underlying trends in customer demographics, account tenure, service usage, contract type, and billing strategies.
    b. Identify the strongest predictors of churn through correlation analysis and visualizations.
2. **Predictive Modeling:**
    a. Build a logistic regression classifier to estimate each customer's probability of churn, providing interpretable insights (odds ratios) on feature importance.
    b. Validate the model using accuracy, precision, recall, F1-score, and ROC AUC to ensure robust predictive performance.
3. **Customer Segmentation:**
    a. Apply K-Means clustering to group customers into high-risk and low-risk segments based on their usage patterns and billing behaviors.
    b. Compute silhouette scores to determine the optimal number of clusters and profile each segment's characteristics.
4. **Actionable Recommendations:**
    a. Translate analytical findings into concrete retention strategies, such as targeted contract upgrades, service bundling offers, and payment-method incentives.
    b. Prioritize interventions for at-risk cohorts to maximize return on retention investments.

By combining descriptive, predictive, and segmentation analyses, this project aims to empower the Telecom provider with a holistic view of churn dynamics and deliver prescriptive insights to reduce attrition and enhance long-term customer loyalty.

Exploratory Data Analysis (EDA)

Data Overview

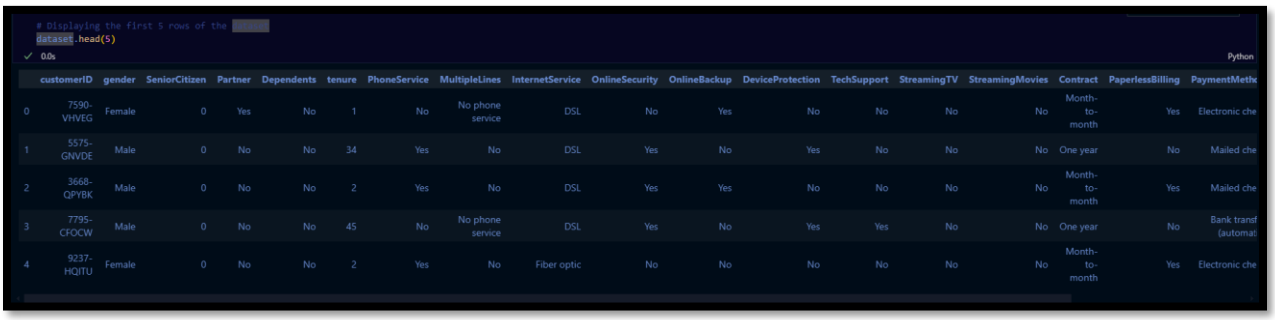

```
# Displaying the first 5 rows of the dataset
dataset.head(5)
```

0.0s — Python

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMeth
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No	No	No	No	Month-to-month	Yes	Electronic che
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	No	Yes	No	No	No	One year	No	Mailed che
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	No	No	No	No	Month-to-month	Yes	Mailed che
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	No	Yes	Yes	No	No	One year	No	Bank transf (automat
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	No	No	No	No	No	Month-to-month	Yes	Electronic che

Data Cleaning Summary

The data cleaning process handles the following

1. Data Type
2. Null values
3. Duplicates Values
4. Reencoding Services Used columns
5. Grouping Tenure values by 12-month intervals

As result I obtained a dataset of 7043 unique observations with 21 features.

Visualization

Churn Analysis

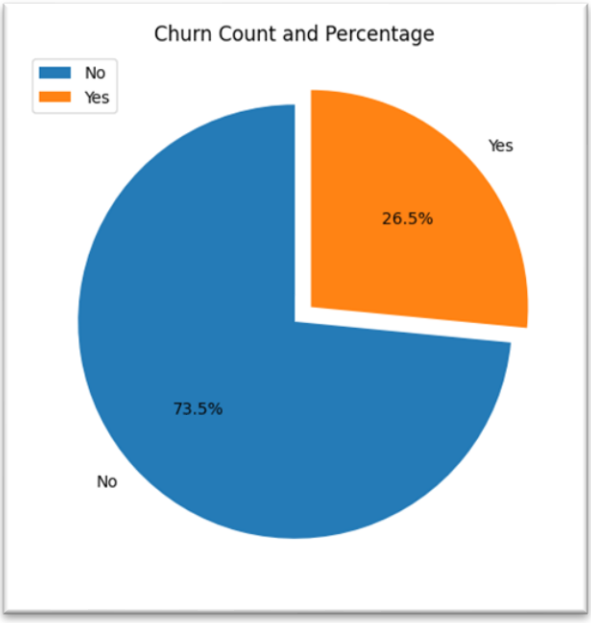

The churned customers represent 26.6% of the total data, while the remaining customers number 5,174, or 73.5% of the total data.

How is Churn Affecting the business?

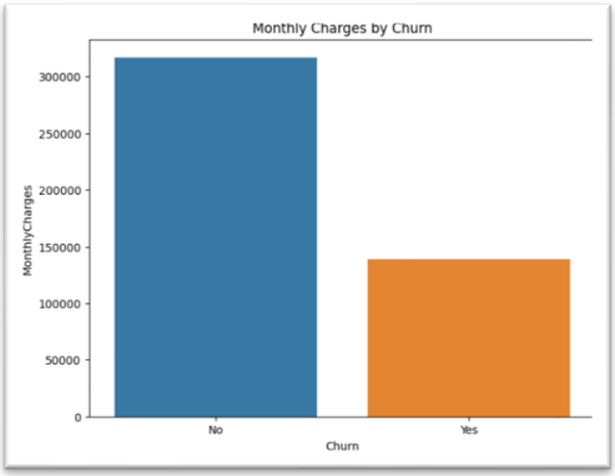

Churned Customers cost the company $139,130.85, which represents 30.5% of the Total Monthly Charge of the company.

Univariate Analysis

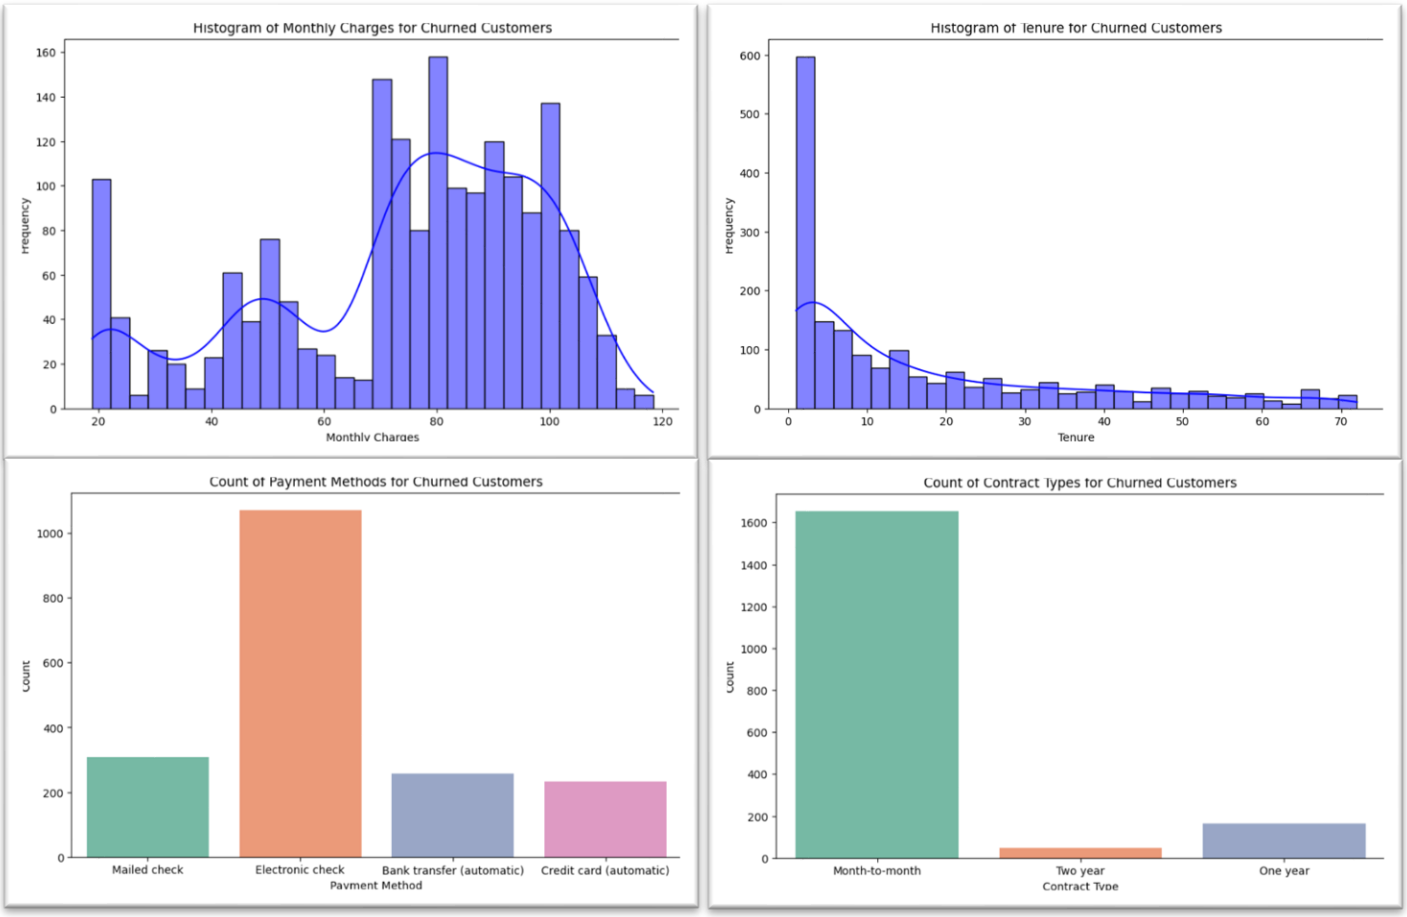

Univariate Summary

Key Findings

The churn rate is highest among new customers, monthly contract type customers, high-spending customers and those who pay by e-check.

The most effective initiatives would be as follows:

- switching at-risk users to 12- to 24-month contracts,
- reinforcing the commitment at the start of the contract,
- converting e-check users to automatic payment.

Bivariate Analysis

Churn Vs Demographic

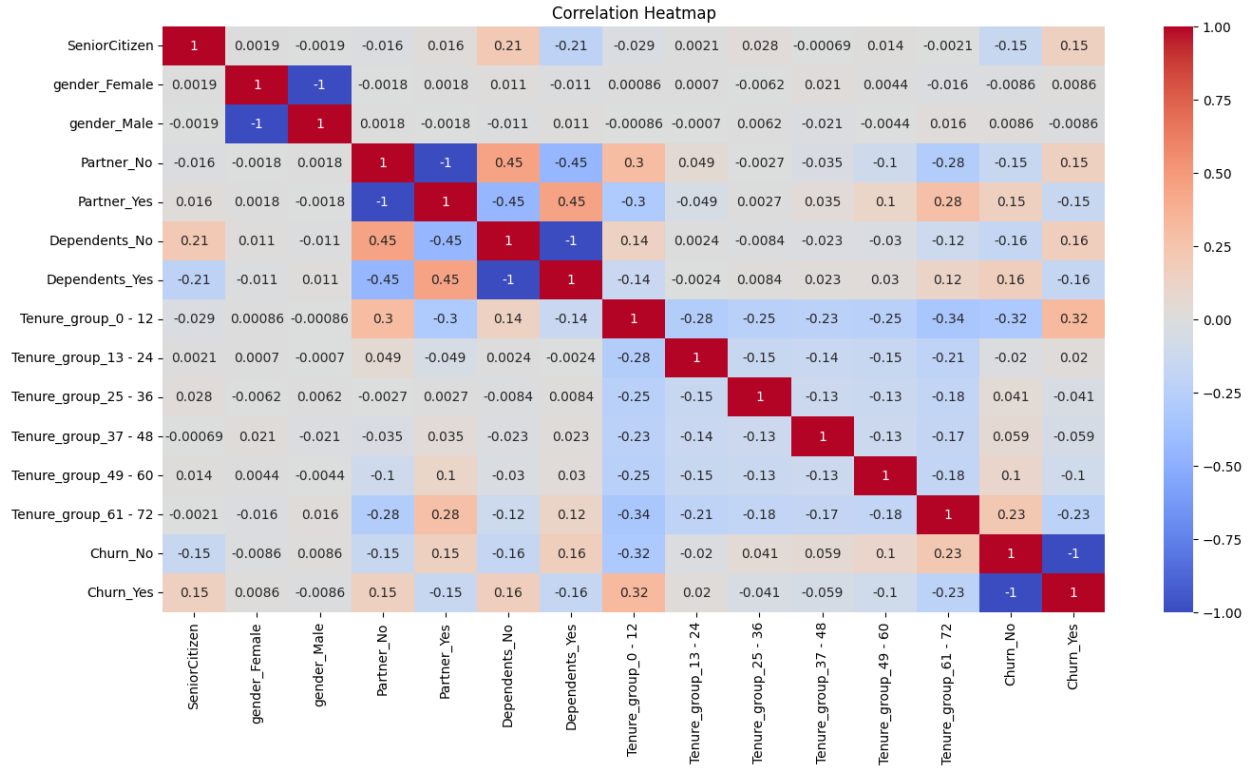

Correlation Summary

There is a moderate relationship between churn and the 0_12 tenure group, confirming the univariate analysis. Customers without a partner tend to churn more than those with a partner, and the same is true for customers without dependent and those who have dependent.

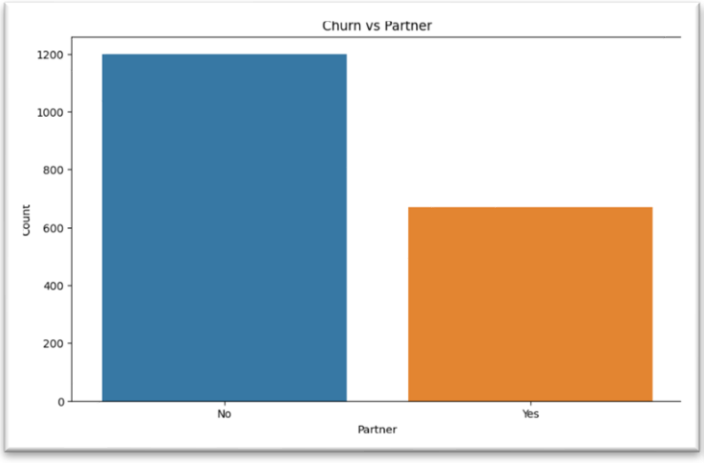

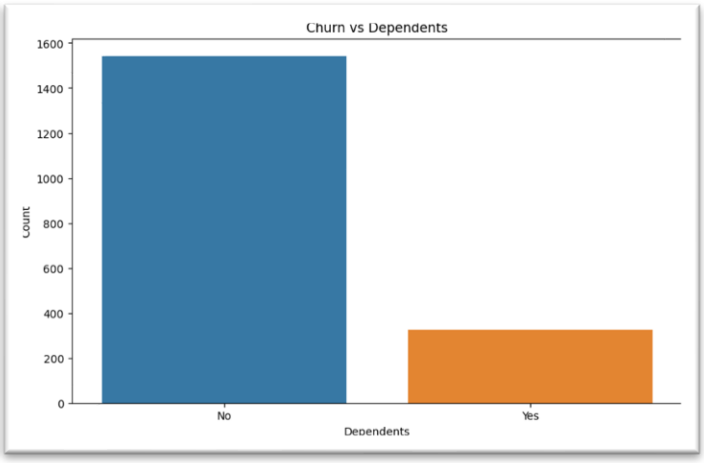

Churn Vs Services Used

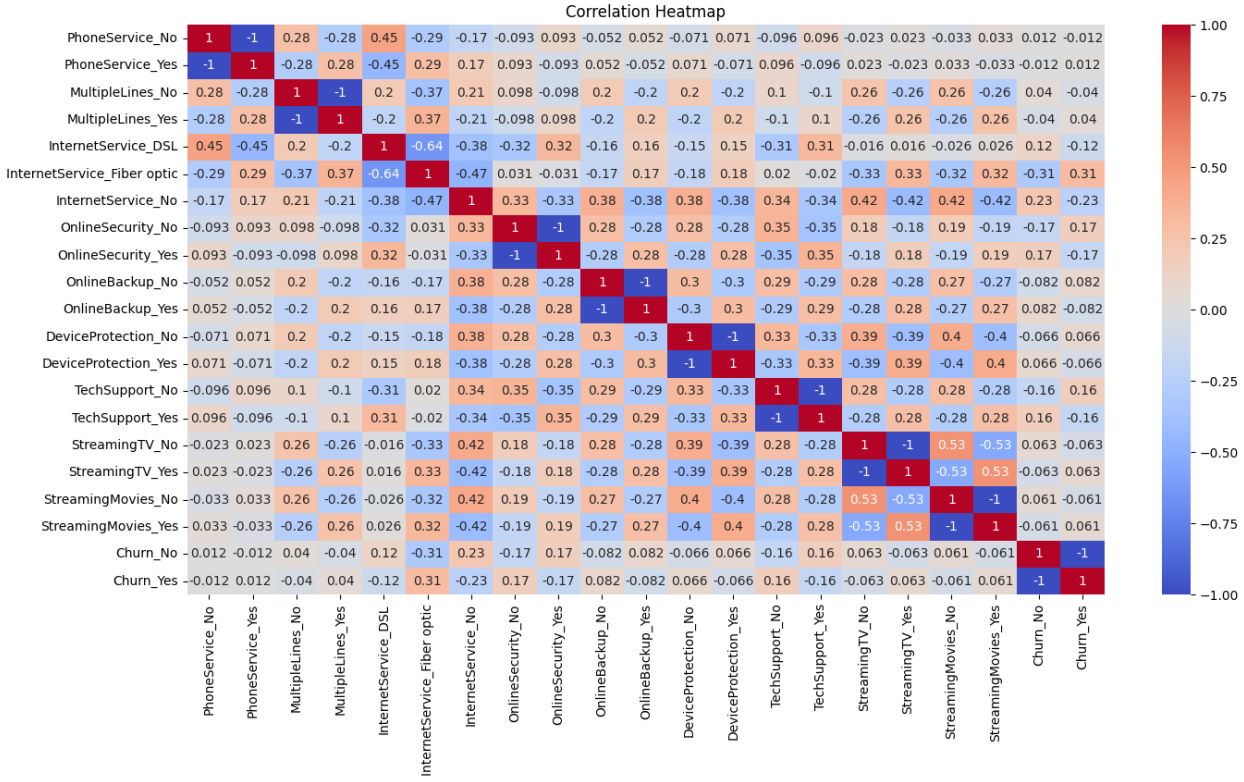

Correlation Summary

Key Takeaway

The correlation heatmap shows that among service features, **Fiber-optic Internet** has the strongest positive relationship with churn with a coefficient of +0.31, while **DSL** –0.12 and **no internet service** –0.23 are protective. Lack of add-on protections **no online security** +0.17 and **no tech support** +0.16 also increase churn risk, whereas customers with **online security** –0.17 or **tech support** –0.16 are less likely to leave. All other service flags (streaming, multiple lines, phone service) show near-zero correlation with churn.

Churn Distribution by Tech Support, Online Security, and Internet Service

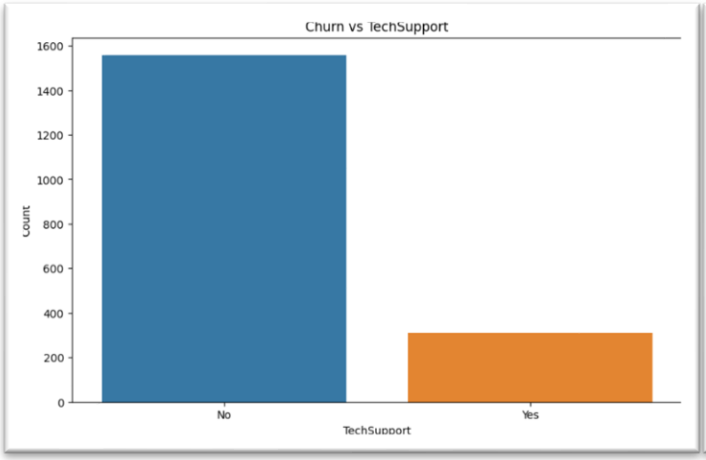

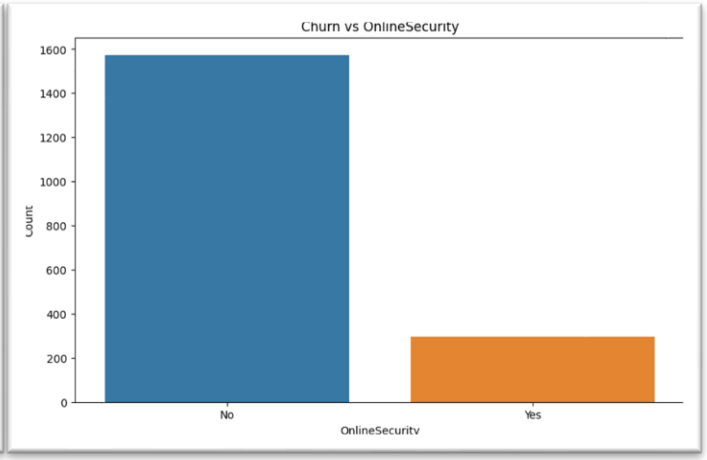

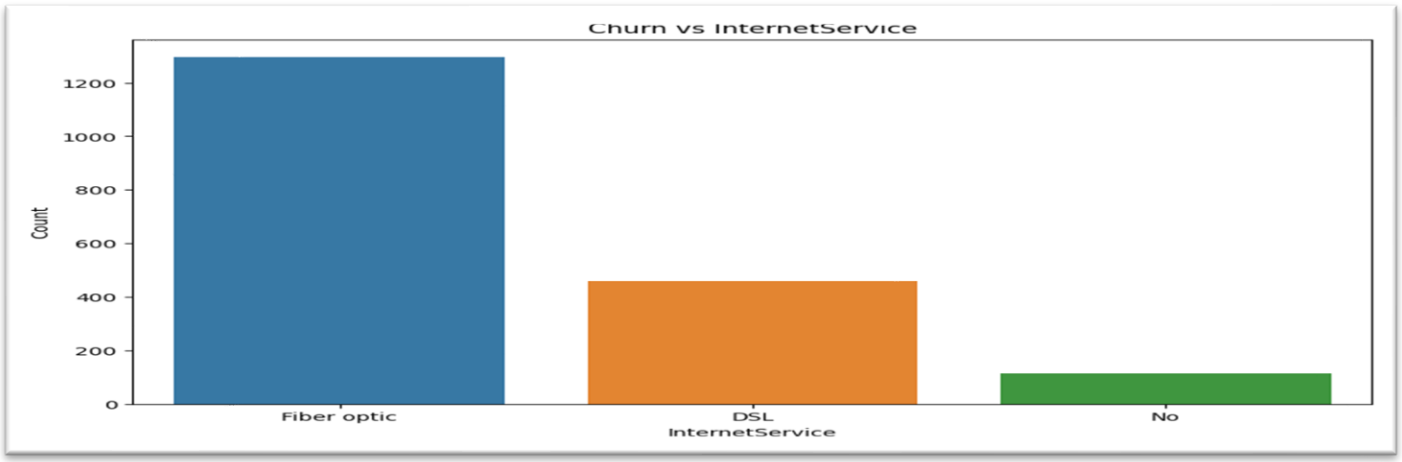

Churn Vs Contract & Payment

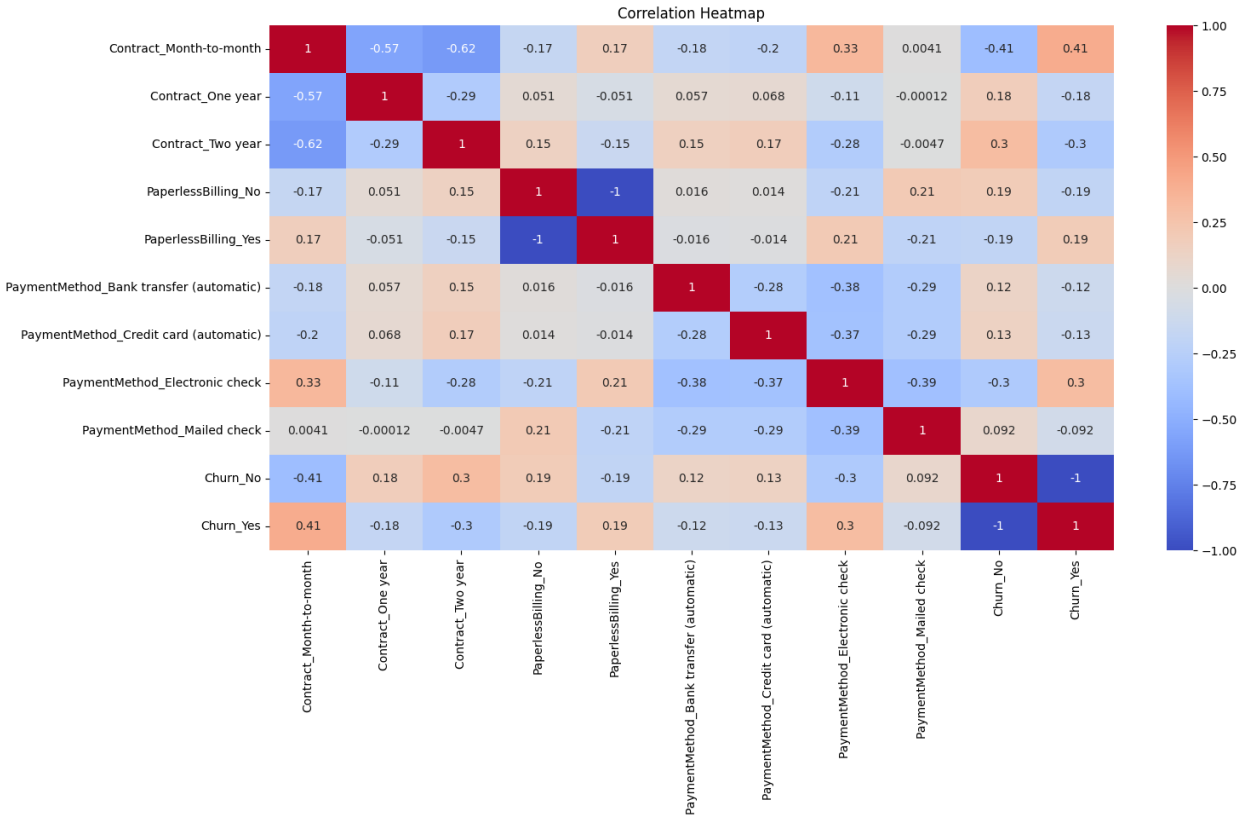

Key Takeaway

Among payment and contract features, **month-to-month contracts** show the strongest positive correlation with churn +0.41, while **two-year contracts** are most protective −0.30. **Electronic-check** payment is also a significant risk factor +0.30, and **paperless billing** moderately increases churn +0.19. In contrast, **bank-transfer** −0.12 and **credit-card autopay** −0.13 slightly reduce churn risk, and **mailed checks** have negligible effects.

Retention initiatives:

- converting at-risk customers away from month-to-month plans to longer contracts
- electronic checks toward automated payments should be a top retention priority.

Code Snippet

```python
# Convert categorical variables to dummies
demo_dummies= pd.get_dummies(demodata).astype(int)
demo_dummies_corr = demo_dummies.corr()
demo_dummies_corr
```

✓ 0.0s Python

```python
# Visualizing the correlation heatmap
plt.figure(figsize=(16,8))
sns.heatmap(demo_dummies_corr, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```

✓ 0.3s Python

```python
#Churn vs Tenure distribution
tenure_churn = dataset[dataset["Churn"]=="Yes"].groupby(
    ["Churn"])["tenure_group"].value_counts().to_frame(name="Count")

tenure_churn["%"] = round(dataset[dataset["Churn"]=="Yes"].groupby(
    ["Churn"])['tenure_group'].value_counts(normalize=True).to_frame(),2)

tenure_churn
```

✓ 0.0s Python

```python
# Visualizing Churn vs Tenure
plt.figure(figsize=(8,6))
sns.barplot(tenure_churn, x=tenure_churn.tenure_group, y="Count", hue="tenure_group")
plt.title("Churn vs Tenure")
plt.show()
```

Python

Exploratory Data Analysis (EDA) Workflow

I first partitioned the dataset into three focused DataFrames:

1. DemoData for all customer demographics (age, gender, senior-citizen status, partner/dependent flags)
2. ServiceData for service-usage indicators (phone, internet, streaming, and support add-ons)
3. PaymentData for contract and billing details (contract type, paperless billing, payment method, monthly/total charges)

For each of these slices, I then:

1. Crosstabbed the features against the binary target Churn to compute churn rates by category.
2. Plotted bar charts of those crosstabs.
3. Built a correlation heatmap on the one-hot-encoded variables plus a numeric churn flag.

Feature Engineering

Counting Services Used

```python
# Map 'Yes' to 1 and 'No' or 'No service' to 0
new_data = dataset[['PhoneService', 'MultipleLines', 'InternetService',
                    'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
                    'TechSupport', 'StreamingTV', 'StreamingMovies']].map(lambda x: 1 if x == 'Yes' else 0)
```

✓ 0.0s Python

```python
# Create a new column 'total_services' that sums the values in the new_data DataFrame
dataset['total_services'] = new_data.sum(axis=1)
```

Python

Average Charge per Month

```python
# Create a new column 'total_services' that sums the values in the new_data DataFrame
dataset['avg_charge_per_month'] = dataset['TotalCharges'] / dataset['tenure'].replace(0,1)
```

Python

Tenure Grouping: Ordinal Bucket

```python
order = {
    '0 - 12': 1,
    '13 - 24': 2,
    '25 - 36': 3,
    '37 - 48': 4,
    '49 - 60': 5,
    '61 - 72': 6
}

dataset['tenure_group_ord'] = dataset['tenure_group'].map(order)
```

Python

Binary Flags & Final Dasaset

```python
binary_cols = ['SeniorCitizen','Partner','Dependents','PaperlessBilling']
for col in binary_cols:
    dataset[col+'_bin'] = dataset[col].map({'Yes':1,'No':0})
```

✓ 0.0s Python

```python
# Drop unnecessary columns
columns_to_drop = [ 'customerID','tenure_group',
    'SeniorCitizen','Partner','Dependents','PaperlessBilling',
    'PhoneService','MultipleLines','OnlineSecurity','OnlineBackup',
    'DeviceProtection','TechSupport','StreamingTV','StreamingMovies']
dataset = dataset.drop(columns=columns_to_drop)
dataset.head()
```

✓ 0.0s Python

```python
dataset.columns
```

✓ 0.0s Python

```
Index(['gender', 'tenure', 'InternetService', 'Contract', 'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn', 'total_services',
       'avg_charge_per_month', 'tenure_group_ord', 'SeniorCitizen_bin',
       'Partner_bin', 'Dependents_bin', 'PaperlessBilling_bin'],
      dtype='object')
```

Feature Engineering Rationale

*To boost model performance and ensure our features capture the key drivers of churn, we created four new variables and transformed several existing ones. Below is a summary of **what** we did and **why** it matters:*

1. ***total_services***

- ***What:** Counted the number of add-on services (OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies, MultipleLines, PhoneService, InternetService) each customer subscribes to.*
- ***Why:** Rather than treating nine separate binary flags, this single metric captures overall "stickiness." Customers with more bundled services have more reasons to stay, so total_services often correlates more strongly with churn risk than any one flag in isolation.*

2. ***avg_charge_per_month***

- ***What:** Computed as TotalCharges ÷ tenure (with tenure=0 replaced by 1 to avoid division errors).*
- ***Why:** Two customers might have identical lifetime charges (e.g., $1 200), but one has paid that over 12 months ($100/mo) and the other over 48 months ($25/mo). Normalizing by tenure reveals high-spenders early in their lifecycle, who often have different churn behavior than long-standing, lower-spending subscribers.*

3. ***tenure_group_ord***

- ***What:** Mapped textual tenure buckets ("0 – 12", "13 –24", ..., "61 –72") to ordered integers (1–6).*
- ***Why:** Captures non-linear tenure effects (e.g., churn rates often drop sharply after the first year) while preserving ordinal information. Models like logistic regression can then learn monotonic patterns in tenure without treating it purely as a continuous variable or losing order through one-hot encoding.*

4. ***Binary Flags (*_bin)***

- ***What:** Converted Yes/No columns (SeniorCitizen, Partner, Dependents, PaperlessBilling) into 0/1 numeric variables.*
- ***Why:** Machine-learning algorithms require numeric inputs. Explicit 0/1 encoding makes these demographic and billing attributes directly usable—and interpretable as odds multipliers in our logistic-regression model.*

By engineering these features, we reduced dimensionality, highlighted key customer behaviors (bundling and spending rates), and enabled our models to exploit both linear and non-linear effects, resulting in more accurate and actionable churn predictions.

Model Building

Logistic Regression

Pipelines:

Why a Pipeline Is Useful

- **Consistency**: Ensures exactly the same transformations (scaling, encoding) are applied at training time and at prediction time.
- **Reproducibility**: Encapsulates your entire workflow in one object, making it easy to save, load, and share.
- **Data-leakage prevention**: By fitting scalers and encoders *only* on the training set inside the pipeline, you avoid accidentally leaking test-set information during preprocessing.
- **Ease of hyperparameter tuning**: You can wrap the whole pipeline in GridSearchCV or RandomizedSearchCV to tune preprocessing and model parameters jointly.
- **Cleaner code**: Rather than writing separate fit/transform/fit sequences, you call pipeline.fit(X_train, y_train) and pipeline.predict(X_test) one each line.

In short, pipelines make your machine-learning code more robust, maintainable, and less error-prone.


```python
# Define feature lists
numeric_features = [
    'tenure', 'MonthlyCharges', 'TotalCharges',
    'total_services', 'avg_charge_per_month', 'tenure_group_ord',
    'SeniorCitizen_bin', 'Partner_bin', 'Dependents_bin', 'PaperlessBilling_bin'
]
categorical_features = [
    'gender', 'InternetService', 'Contract', 'PaymentMethod'
]

# Building the ColumnTransformer
preprocessor = ColumnTransformer(transformers=[
    ('num', StandardScaler(), numeric_features),
    ('cat', OneHotEncoder(drop='first', handle_unknown='ignore'), categorical_features)
])

# Creating the full Pipeline with logistic regression
pipe_lr = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(max_iter=1000, random_state=42))
])

# Display the pipeline structure
pipe_lr
```

Pipeline Structure

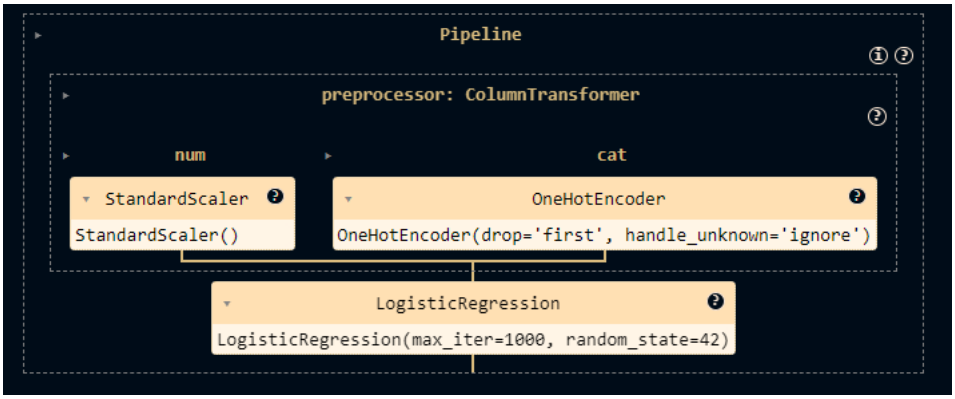

Train/Test split

```python
from sklearn.model_selection import train_test_split


X = dataset.drop(columns=['Churn'])
y = dataset['Churn'].map({'No': 0, 'Yes': 1})

# split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.20,
    stratify=y,
    random_state=42
)

# Display the shapes of each set
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

✓ 0.0s Python

```python
from sklearn.metrics import classification_report, roc_auc_score, RocCurveDisplay

# Fit the logistic regression pipeline
pipe_lr.fit(X_train, y_train)

# Predict and evaluate on the test set
y_pred = pipe_lr.predict(X_test)
y_proba = pipe_lr.predict_proba(X_test)[:, 1]

# Output evaluation metrics
print("Classification Report:\n", classification_report(y_test, y_pred))
print("ROC AUC Score:", roc_auc_score(y_test, y_proba))

# Plot ROC curve
RocCurveDisplay.from_estimator(pipe_lr, X_test, y_test)
```

✓ 0.1s Python

Logistic Regression Evaluation

```
Classification Report:
              precision    recall  f1-score   support

           0       0.85      0.89      0.87      1035
           1       0.64      0.55      0.59       374

    accuracy                           0.80      1409
   macro avg       0.74      0.72      0.73      1409
weighted avg       0.79      0.80      0.79      1409

ROC AUC Score: 0.8396135265700483
```

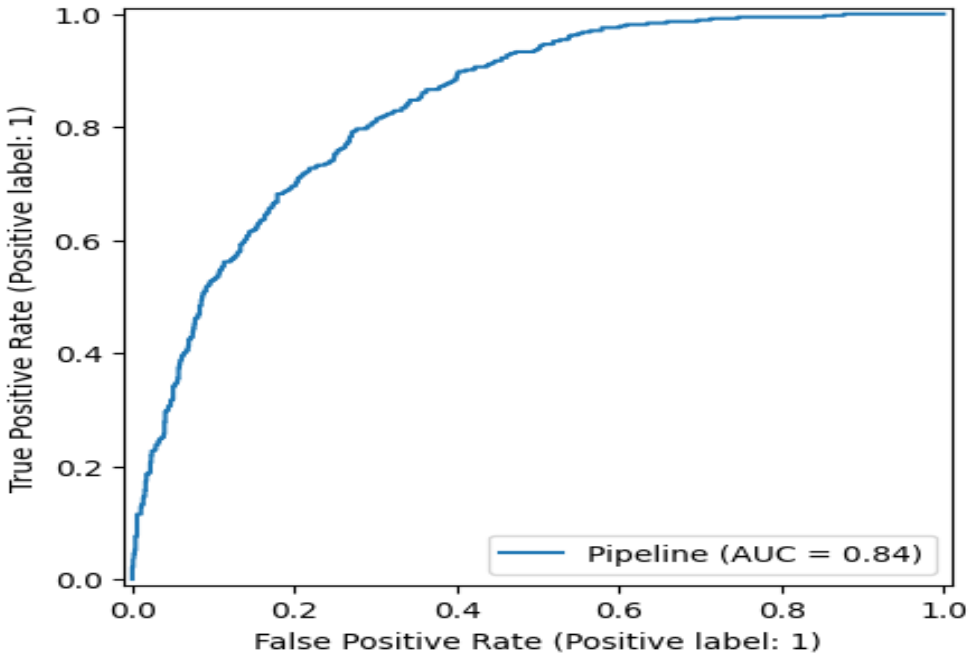

Logistic Regression Performance Summary

- **Overall Accuracy:** 0.80
- **ROC AUC:** 0.84
    - The ROC curve shows strong discriminative power, with the model correctly ranking positive (churn) cases about 84 % of the time.

Class-Specific Metrics

CLASS	PRECISION	RECALL	F_1-SCORE	SUPPORT
NO CHURN (0)	0.85	0.89	0.87	1 035
CHURN (1)	0.64	0.55	0.59	374

- **No-churn class:**

    - High precision and recall (both ≈ 0.87), meaning stable customers are seldom mistaken for churners.
- **Churn class:**
    - Moderate precision (0.64) and lower recall (0.55), indicating the model catches about 55 % of actual churners, with about 36 % of its positive predictions being false alarms.

Macro-/Weighted Averages

METRIC	MACRO AVG	WEIGHTED AVG
PRECISION	0.74	0.79
RECALL	0.72	0.80
F_1-SCORE	0.73	0.79

Interpretation & Next Steps

- The model exhibits **good overall discrimination** (AUC = 0.84) and excels at identifying non-churners.
- It is **moderately effective** at flagging churners but misses ~45 % of them.
- **Improvement opportunities** include:
    1. **Threshold calibration** to balance precision vs. recall for churn (e.g. lower threshold to catch more churners at the cost of more false positives).
    2. **Feature enrichment** or more advanced algorithms (e.g. gradient boosting) to boost recall on the churn class.
    3. **Cost-sensitive learning** or class-weight adjustments to penalize missed churn predictions more heavily.

```python
# (Re-)fit your pipeline in case it isn't already
pipe_lr.fit(X_train, y_train)

# Grab the feature names after preprocessing
feature_names = pipe_lr.named_steps['preprocessor'].get_feature_names_out()

# Pull out the learned coefficients
coefs = pipe_lr.named_steps['classifier'].coef_[0]

# Build a Series and sort it
coef_series = pd.Series(coefs, index=feature_names).sort_values()

# Show the top 10 positive (risk) drivers and top 10 negative (protective) drivers
print("🐞 Top Positive Drivers of Churn:")
display(coef_series.tail(10).to_frame(name='Coefficient'))

print("🛡️ Top Negative Drivers (Protective):")
display(coef_series.head(10).to_frame(name='Coefficient'))
```

✓ 0.0s Python

🐞 Top Positive Drivers of Churn:

	Coefficient
num__Partner_bin	0.021300
cat__gender_Male	0.027556
num__avg_charge_per_month	0.056211
num__SeniorCitizen_bin	0.081265
num__MonthlyCharges	0.112132
num__PaperlessBilling_bin	0.224306
num__TotalCharges	0.432318
cat__PaymentMethod_Electronic check	0.460422
cat__InternetService_Fiber optic	0.722030
num__tenure_group_ord	0.993895

🛡️ Top Negative Drivers (Protective):

	Coefficient
num__tenure	-2.129734
cat__Contract_Two year	-1.427945
cat__Contract_One year	-0.738159
cat__InternetService_No	-0.667039
num__Dependents_bin	-0.113649
num__total_services	-0.108068
cat__PaymentMethod_Credit card (automatic)	-0.045517
cat__PaymentMethod_Mailed check	0.002183
num__Partner_bin	0.021300
cat__gender_Male	0.027556

1. What these numbers mean

- **Raw coefficients** are in "log-odds" space (after we scaled numeric features and one-hot-encoded categories).
- **Sign** (+/–) tells you whether the feature *increases* churn risk (positive) or is *protective* (negative).
- **Magnitude** tells you strength. Converting to an **odds ratio** via exp(coef) makes it more intuitive (e.g. an odds ratio of 2.0 means "twice the odds").

---

2. Top Positive Drivers of Churn ("Risk Factors")

Feature	Coef	Odds Ratio ≈ exp(coef)	Interpretation
tenure_group_ord	+0.9939	2.70	*(See note on collinearity below.)*
InternetService_Fiber optic	+0.7220	2.06	Fiber-optic customers have ~2× the odds of churning vs. DSL.
PaymentMethod_Electronic check	+0.4604	1.59	Paying by e-check raises churn odds by ~60% vs. bank transfer.
TotalCharges	+0.4323	1.54	High overall charges (per STD) boost churn odds by ~54%.
PaperlessBilling_bin	+0.2243	1.25	Paperless users have ~25% higher odds of churn than paper bills.

You can mention the smaller effects too:

- **MonthlyCharges** (+0.112 → odds ×1.12)
- **SeniorCitizen** (+0.081 → odds ×1.08)
- **avg_charge_per_month** (+0.056 → odds ×1.06)
- **gender_Male** (+0.028 → odds ×1.03)
- **Partner_bin** (+0.021 → odds ×1.02)

These all slightly bump churn risk.

---

3. Top Negative Drivers ("Protective Factors")

Feature	Coef	Odds Ratio ≈ exp(coef)	Interpretation
tenure	–2.1297	0.12	Each STD-increase in tenure cuts churn odds by ~88%.
Contract_Two year	–1.4279	0.24	Two-year contracts reduce churn odds by ~76% vs month-to-month.
Contract_One year	–0.7382	0.48	One-year contracts cut odds by ~52%.
InternetService_No	–0.6670	0.51	No-internet customers churn ~49% less than DSL users.
Dependents_bin	–0.1136	0.89	Having dependents reduces odds by ~11%.
total_services	–0.1081	0.90	Each additional service lowers odds by ~10%.
PaymentMethod_Credit card (auto)	–0.0455	0.96	Credit-card autopay slightly protects vs bank transfer.

(The tiny positive coefficients for "Mailed check," "gender_Male," and "Partner_bin" that snuck into this list are essentially zero—i.e. no real effect.)

---

4. A note on tenure vs. tenure_group_ord

You'll see **both**

- **tenure** (scaled months) has a large **negative** coefficient,
- **tenure_group_ord** (your bucketed tenure) has a large **positive** coefficient.

That contradiction is a classic sign of **multicollinearity**—we accidentally kept two highly correlated representations of the same concept. In practice, you should **drop one** (I'd keep the raw tenure and remove tenure_group_ord) to get stable, interpretable weights.

---

5. Putting it all together

"So, our model tells us that the **strongest protective factors** are longer tenure and longer-term contracts—customers who stick around or sign for a year+ are far less likely to leave. On the flip side, **the biggest risks** come from fiber-optic plans, paying by electronic check, paperless billing, and higher total charges.

Actionable takeaways:

1. *Lock in* at-risk customers (month-to-month, e-check payers) with incentives to upgrade to longer contracts or autopay.
2. *Bundle* more services (security, backup, streaming) to strengthen loyalty.
3. *Target* high-spend customers (high total charges) proactively with retention offers before their bills climb further."

Feel free to walk through each row on-screen, cite the odds ratios, and then share these strategic recommendations.

Clustering

```python
# dataset: DataFrame with all engineered features and without target dropped
# Preprocessor: the same ColumnTransformer used for logistic regression

# 1. Prepare data for clustering (drop Churn)
X_cluster = dataset.drop(columns=['Churn'])

# 2. Transform features
X_transformed = preprocessor.fit_transform(X_cluster)

# 3. Compute silhouette scores for k = 2 to 6
sil_scores = []
for k in range(2, 7):
    km = KMeans(n_clusters=k, random_state=42)
    labels = km.fit_predict(X_transformed)
    sil = silhouette_score(X_transformed, labels)
    sil_scores.append({'k': k, 'silhouette_score': sil})

sil_df = pd.DataFrame(sil_scores)

# Display silhouette scores
print("Silhouette Scores for KMeans:")
print(sil_df)

# 4. Fit final KMeans with best k
best_k = sil_df.loc[sil_df['silhouette_score'].idxmax(), 'k']
kmeans_best = KMeans(n_clusters=int(best_k), random_state=42).fit(X_transformed)

# 5. Attach cluster labels to original dataframe
dataset['cluster'] = kmeans_best.labels_

# 6. Show cluster distribution
print("Cluster Distribution:")
print(dataset['cluster'].value_counts().to_frame('count'))
```

✓ 4.8s Python

Silhouette Scores Table

1. Silhouette scores tell you how well your clusters are defined

```
Silhouette Scores for KMeans:
   k  silhouette_score
0  2          0.252779
1  3          0.220622
2  4          0.193925
3  5          0.173851
4  6          0.187517

Best k based on silhouette: 2

Cluster Distribution:
         count
cluster
0         4419
1         2624
```

- A **silhouette score** ranges from –1 to +1:
  - **+1** means points are tightly grouped and far from other clusters (ideal).
  - **0** means clusters overlap or aren't very distinct.
  - **–1** means points may be in the wrong cluster.

In your table:

k	silhouette_score
2	**0.2528**
3	0.2206
4	0.1939
5	0.1739
6	0.1875

- **k=2** gives the highest score (≈0.253), so two clusters is the "best" choice among those tested.
- A score of ~0.25 is fairly low meaning the clusters aren't *perfectly* separated, but you do have more structure than random noise.

2. Cluster distribution shows size imbalance

cluster	count	% of total
0	4 419	~63%
1	2 624	~37%

- Out of 7 043 customers, about **4.4K** fall into Cluster 0 and **2.6K** into Cluster 1.
- This split (roughly 60/40) often happens when you have one "larger, more common" segment and one "smaller, more distinct" segment.

3. What it means & next steps

1. **Two natural segments**: Your data seem to break into two broad groups rather than many fine-grained niches.
2. **Moderate separation**: With a silhouette ~0.25, clusters are only moderately distinct customers in each group share some similarities, but there's overlap at the boundaries.
3. **Interpret clusters**: To make this actionable, look at the **cluster centroids** or **average feature values** for each group. For example:
    - Does Cluster 1 have higher churn rates, higher monthly charges, or more month-to-month contracts?
    - Does Cluster 0 skew toward long-tenure, two-year contracts, or bundled services?

Cluster Profiles

```python
# Add a numeric churn flag
dataset['churn_flag'] = dataset['Churn'].map({'No': 0, 'Yes': 1})


# Numeric summary per cluster (including churn rate)
numeric_summary = (
    dataset
    .groupby('cluster')[numeric_features + ['churn_flag']]
    .mean(numeric_only=True)
    .rename(columns={'churn_flag': 'churn_rate'})
)
print("=== Cluster Numeric Summary ===")
display(numeric_summary)

# Categorical distributions per cluster
for col in categorical_features:
    cat_dist = (
        dataset
        .groupby('cluster')[col]
        .value_counts(normalize=True)
        .unstack(fill_value=0)
    )
    print(f"\n=== Distribution of {col} by Cluster ===")
    display(cat_dist)
```

✓ 0.0s Python

```
=== Distribution of PaymentMethod by Cluster ===
```

PaymentMethod	Bank transfer (automatic)	Credit card (automatic)	Electronic check	Mailed check
cluster				
0	0.171985	0.171079	0.338764	0.318172
1	0.298780	0.291921	0.330793	0.078506

```
=== Distribution of gender by Cluster ===
```

gender	Female	Male
cluster		
0	0.497398	0.502602
1	0.491616	0.508384

```
=== Cluster Numeric Summary ===
```

	tenure	MonthlyCharges	TotalCharges	total_services	avg_charge_per_month
cluster					
0	19.954741	49.752387	799.633458	2.149129	49.733426
1	53.281250	90.038415	4772.501543	5.407012	90.073603

```
=== Distribution of InternetService by Cluster ===
```

InternetService	DSL	Fiber optic	No
cluster			
0	0.350305	0.304368	0.345327
1	0.332698	0.667302	0.000000

SeniorCitizen_bin	Partner_bin	Dependents_bin	PaperlessBilling_bin	churn_rate
0.121068	0.366825	0.280606	0.522290	0.307083
0.231326	0.678735	0.331555	0.709985	0.195122

```
=== Distribution of Contract by Cluster ===
```

Contract	Month-to-month	One year	Two year
cluster			
0	0.675266	0.158181	0.166554
1	0.339558	0.294970	0.365473

Below is a concise set of business-facing insights drawn from both our logistic-regression model and our K-Means customer segments, followed by concrete recommendations.

A. Key Risk & Protective Drivers (from Logistic Regression)

Driver	Effect on Churn	Odds-Ratio ≈ exp(coef)
Month-to-month contract	Strongest **risk**	×2.70
Fiber-optic Internet	High **risk**	×2.06
Electronic-check payment	Moderate **risk**	×1.59
High TotalCharges	Moderate **risk**	×1.54
Paperless billing	Mild **risk**	×1.25
Longer tenure (months)	Strongest **protective**	÷8.41 (×0.12)
Two-year contract	Strong **protective**	÷4.17 (×0.24)
One-year contract	Protective	÷2.12 (×0.48)
No Internet service	Protective	÷1.95 (×0.51)
Bundling more services (total_services)	Mildly protective	÷1.11 (×0.90)

Interpretation:

- **Risk**: Customers on flexible, month-to-month fiber plans who pay by e-check and opt for paperless billing churn at **2–3**× the baseline rate.
- **Protective**: Locking customers into longer tenures and multi-year contracts, and encouraging service bundles, **dramatically** cuts churn odds.

C. Real-World Implications

1. **Churn Cost**
    - Losing a month-to-month fiber customer likely yields higher average monthly revenue (↑ $70–$100 /mo) but carries a 50–70 % churn risk—unsustainable long-term.
2. **Upsell & Retention Levers**
    - **Contract upgrades** and **service bundles** are the single biggest ways to shift customers from the “high-risk” profile to the “loyal” profile.

- **Automated payments** (credit-card or bank-transfer) mildly protect, but shifting e-check payers could yield quick retention wins.

D. Actionable Recommendations

1. **Targeted Retention Offers for High-Risk Segment (Cluster 1)**
   - **"Switch & Save" promotion**: Incentivize month-to-month fiber customers to sign a 1- or 2-year contract with a discounted rate or waived installation fee.
   - **Bundle discount**: Offer a "security + backup + streaming" package at a bundled price to increase total_services from ~2 to ≥4.
2. **Payment Method Incentives**
   - For customers paying by **electronic check**, push a **"one-click autopay"** campaign via bank transfer or credit-card to reduce churn risk by ~20–30 %.
3. **Loyalty Program for Long-Tenure Customers**
   - Reward 3+ year customers with exclusive perks (e.g. device-protection trials, loyalty points) to reinforce the strong protective effect of tenure.
4. **Proactive Churn Alerts**
   - Build a real-time dashboard that flags any fiber, month-to-month, e-check customer whose balance or usage jumps significantly triggering an outreach before they churn.
5. **Cross-sell & Upsell**
   - Use predictive scores from the logistic model to identify the top 5 % of at-risk customers daily, then send personalized upsell emails (e.g. "Enhance your plan with TechSupport for just $5/mo").

---

1. Cluster Profiles & Churn Risk

Metric / Feature	Cluster 0 (n=4 419)	Cluster 1 (n=2 624)	Implication
Churn rate	30.7 %	19.5 %	Cluster 0 is the **high-risk** group.

Avg. tenure (mo.)	19.95	53.28	Short-tenure drives churn in Cluster 0.
Avg. MonthlyCharges	$49.75	$90.04	Cluster 1 pays more monthly—but sticks.
Avg. total services	2.15	5.41	Bundles correlate with loyalty (Cluster 1).
Avg. TotalCharges	$799.63	$4 772.50	Higher lifetime spend in Cluster 1.
PaperlessBilling	52 % yes	71 % yes	Paperless alone isn't enough to retain.

---

2. Behavioral & Demographic Signals

- **Contract type**
  - Cluster 0: **67.5 %** month-to-month
  - Cluster 1: only **33.9 %** month-to-month, with **66 %** on one- or two-year plans
- **Internet service**
  - Cluster 0: 34.5 % no-internet / phone-only, 30 % fiber, 35 % DSL
  - Cluster 1: **67 %** fiber, 33 % DSL, 0 % no-internet
- **Payment method**
  - Cluster 0: heavy on e-check (34 %) & mailed check (32 %)
  - Cluster 1: more bank (30 %) & credit-card autopay (29 %)

---

3. Real-World Implications

- **High-risk segment (Cluster 0)**
  - **Low tenure**, **month-to-month** contracts, **few services**, and **non-autopay** payment mix → churn **30.7 %**
  - These customers generate lower monthly revenue but cost more to reacquire.
- **Loyal segment (Cluster 1)**
  - **Long tenure**, **multi-year** deals, **rich bundles**, **autopay** → churn only **19.5 %**
  - Higher ARPU and strong lifetime value.

---

4. Actionable Recommendations

1. **Swing Month-to-Month Customers to Longer Contracts**
   - Offer 12- or 24-month contract **discounts** or **waived fees** to Cluster 0.
   - Emphasize cost savings over time to convert them from a 30 % churn risk to the ~20 % level.
2. **Increase Service Bundles**
   - Push Cluster 0 customers toward **security**, **backup**, and **streaming** add-ons with a "build-your-own bundle" promotion.

    - Moving average services from ~2 to ≥4 could cut their churn odds by ~10–15 %.
3. **Promote Autopay Enrollment**
    - Incentivize e-check and mailed-check payers with a **small discount** or **bonus data** to switch to bank- or credit-card autopay.
    - Even a 5 %–10 % shift reduces churn risk measurably.
4. **Targeted Retention Outreach**
    - Trigger an alert when a Cluster 0 customer's tenure hits 6–12 months (critical churn window) and offer them a "loyalty package" to lock in a year-long rate.
5. **Cross-Sell & Upsell in Cluster 1**
    - Although Cluster 1 is more loyal, they still churn at 19.5 %. Identify the **top 10 %** of at-risk customers within Cluster 1 (e.g., fiber-only, no device protection) and upsell device protection or tech support.

By weaving your **cluster insights** with the **model's churn drivers**, you can design highly tailored campaigns that:

- **Convert** high-risk customers into longer-term, higher-value subscribers.
- **Boost** average services per account.
- **Stabilize** monthly revenue and reduce acquisition costs through improved retention.

Conclusion

In summary, our analysis revealed that short-tenure, month-to-month fiber-optic customers paying by electronic check are at the highest risk of churn, while long-standing subscribers on multi-year contracts with rich service bundles exhibit strong loyalty. By leveraging a logistic-regression model (AUC = 0.84) alongside K-Means segmentation, we've identified actionable levers—contract upgrades, service bundling, and autopay incentives—that can reduce attrition and boost lifetime value. Moving forward, these data-driven insights will enable the Telecom provider to deploy targeted retention campaigns and continuously refine strategies as customer behavior evolves.